

## Java 8

### **1. What is Functional Interface?**

A functional interface is an interface that has exactly one abstract method. This single abstract method represents the functional behaviour of the interface. Functional interfaces are the foundation of lambda expressions in Java. They can also contain default methods and static methods without breaking the rule of one abstract method. The annotation `FunctionalInterface` is optional but recommended because it gives compile time error if more than one abstract method is added.

### **2. What are the functional interfaces that were already present before Java 8?**

`Runnable` represents a task that runs in a separate thread. It has only one method `run` and does not return any value.

`Callable` is similar to `Runnable` but it returns a value and can throw checked exceptions. It is mainly used with `ExecutorService`.

`Comparator` is used to define custom sorting logic. It has one abstract method `compare` which compares two objects.

`Comparable` is used to define natural ordering of objects. It has one method `compareTo` which compares the current object with another object.

### **3. Can we extend one functional interface from another functional interface?**

Yes one functional interface can extend another functional interface as long as the child interface does not introduce an additional abstract method. If it adds another abstract method then it will no longer be a functional interface. Default methods do not affect this rule.

### **4. What are all the functional interfaces introduced in Java 8?**

**\*\*Predicate\*\***

`Predicate` is used to test a condition and return true or false. It is mostly used in stream filtering.

Method name is `test` which takes one input and returns boolean.

Example using stream

```
```java
List<Integer> numbers = List.of(10, 15, 20, 25);

numbers.stream()
    .filter(n -> n % 2 == 0)
    .forEach(System.out::println);
````
```

Here filter internally calls the `test` method.

**\*\*Function\*\***

Function is used to transform one value into another value. It is commonly used in map operations.

Method name is apply which takes one input and returns one output.

Example using stream

```
```java
List<Integer> numbers = List.of(2, 3, 4);

numbers.stream()
    .map(n -> n * n)
    .forEach(System.out::println);
```

```

Here map internally calls the apply method.

**\*\*Consumer\*\***

Consumer is used to perform an action on the given input. It does not return anything.

Method name is accept which takes one input and returns nothing.

Example using stream

```
```java
List<String> names = List.of("A", "B", "C");

names.stream()
    .forEach(name -> System.out.println(name));
```

```

Here forEach internally calls the accept method.

**\*\*Supplier\*\***

Supplier is used to provide or generate a value. It does not take any input.

Method name is get which returns a value.

Example

```
```java
Supplier<String> supplier = () -> "Java";
System.out.println(supplier.get());
```

```

### \*\*UnaryOperator\*\*

UnaryOperator is a special type of Function where input and output types are same.

Method name is apply inherited from Function.

Example using stream

```
```java
List<Integer> numbers = List.of(1, 2, 3);

numbers.stream()
    .map(n -> n + 1)
    .forEach(System.out::println);
```

```

### \*\*BinaryOperator\*\*

BinaryOperator is a special type of BiFunction where both inputs and output are of same type.

Method name is apply inherited from BiFunction.

Example using stream

```
```java
List<Integer> numbers = List.of(1, 2, 3, 4);

int sum = numbers.stream()
    .reduce(0, (a, b) -> a + b);
```

```

```
System.out.println(sum);
```

```

## **\*\*BiPredicate\*\***

BiPredicate takes two inputs and returns true or false.

Method name is test which takes two inputs and returns boolean.

### Example

```
```java
```

```
BiPredicate<Integer, Integer> bp = (a, b) -> a + b > 10;
```

```
System.out.println(bp.test(5, 7));
```

```
```
```

## **\*\*BiFunction\*\***

BiFunction takes two inputs and returns one output.

Method name is apply which takes two inputs and returns one output.

### Example

```
```java
```

```
BiFunction<String, String, String> bf =
```

```
    (name, city) -> name + " from " + city;
```

```
System.out.println(bf.apply("Ram", "Delhi"));
```

```
```
```

## **\*\*BiConsumer\*\***

BiConsumer takes two inputs and performs an action without returning anything.

Method name is accept which takes two inputs and returns nothing.

### Example using stream

```
```java
```

```
Map<Integer, String> map = Map.of(1, "A", 2, "B");
```

```
map.forEach((key, value) ->  
    System.out.println(key + " " + value));  
...  
...
```

## 5. What is lambda?

Lambda is a short form of writing an implementation of a functional interface. Instead of creating a class or anonymous inner class you directly write the logic. Lambda expression represents the implementation of the single abstract method of a functional interface.

## 6. What are the advantages and disadvantages of using lambda expression?

Advantages are reduced boilerplate code better readability for small logic and support for functional programming style.

Disadvantages are harder debugging stack traces are less clear and complex lambda expressions can reduce readability.

## 7. What is stream?

Stream is a feature introduced in Java 8 that allows us to process collections of data in a functional and declarative way. A stream represents a sequence of elements coming from a data source like a list, set, or array, and it supports operations such as filtering, mapping, sorting, and aggregation. Stream does not store data and does not modify the original collection. It only processes the data.

Streams work using lazy execution, which means operations are not executed until a terminal operation like `forEach` or `collect` is called. This helps improve performance because only required elements are processed. Streams can also be processed sequentially or in parallel. In real applications, streams are used to write clean, readable, and efficient code for data processing.

## 8. What is method reference?

Method reference is a shorter and cleaner way of writing a lambda expression when the lambda only calls an existing method. Instead of writing the full lambda syntax, we directly refer to the method using the class name or object name. This improves code readability and makes the intention very clear. Method reference works with functional interfaces because the referenced method must match the abstract method signature of the functional interface.

In simple terms, if a lambda expression only does one thing and that thing is calling another method, then method reference is a better choice.

There are four types of method references.

First is static method reference. It is used when the method being called is static. For example, when printing elements of a list, instead of writing a lambda, we can write `list.forEach(System.out::println)`. Here `println` is a static method reference and it matches the `Consumer` interface.

Second is instance method reference of a particular object. It is used when the method belongs to a specific object. For example, if we have a `Printer` object and want to print each element, we can write `names.forEach(printer::print)`. The `print` method belongs to the `printer` object.

Third is instance method reference of an arbitrary object of a class. It is used when the method belongs to any object of a particular class. A common example is `names.stream().map(String::toUpperCase)`. Here `toUpperCase` is called on each `String` object in the stream.

Fourth is constructor reference. It is used to refer to a constructor. For example, names.stream().map(Employee::new).toList() creates Employee objects using the constructor.

## 9. Stream methods

filter is an intermediate operation used to select elements based on condition.

map is an intermediate operation used to transform elements.

flatMap is an intermediate operation used to flatten nested structures.

sorted is an intermediate operation used to sort elements.

forEach is a terminal operation used to iterate elements.

collect is a terminal operation used to convert stream into list set or map.

reduce is a terminal operation used to combine elements into a single result.

### filter

filter is used to select elements based on a condition. It internally uses Predicate.

Example

```
List<Integer> numbers = List.of(10, 15, 20, 25);
```

```
numbers.stream()
```

```
    .filter(n -> n > 15)
```

```
    .forEach(System.out::println);
```

### map

map is used to transform each element into another form. It internally uses Function.

Example

```
List<Integer> numbers = List.of(1, 2, 3);
```

```
numbers.stream()
```

```
    .map(n -> n * 2)
```

```
    .forEach(System.out::println);
```

### flatMap

flatMap is used to flatten nested data structures into a single stream.

Example

```
List<List<Integer>> list = List.of(
```

```
    List.of(1, 2),
```

```
    List.of(3, 4)
```

```
);
```

```
list.stream()  
.flatMap(l -> l.stream())  
.forEach(System.out::println);
```

### **sorted**

sorted is used to sort elements in natural or custom order.

#### Example

```
List<Integer> numbers = List.of(30, 10, 20);
```

```
numbers.stream()
```

```
.sorted()  
.forEach(System.out::println);
```

### **forEach**

forEach is a terminal operation used to perform an action on each element. It internally uses Consumer.

#### Example

```
List<String> names = List.of("A", "C", "B");
```

```
names.stream()
```

```
.forEach(name -> System.out.println(name));
```

### **collect**

collect is used to convert a stream into a collection or another data structure.

#### Example

```
List<Integer> numbers = List.of(1, 2, 3);
```

```
List<Integer> result = numbers.stream()
```

```
.map(n -> n * 2)  
.collect(Collectors.toList());
```

```
System.out.println(result);
```

### **reduce**

reduce is used to combine all elements into a single result. It internally uses BinaryOperator.

### Example

```
List<Integer> numbers = List.of(1, 2, 3, 4);
```

```
int sum = numbers.stream()
```

```
.reduce(0, (a, b) -> a + b);
```

```
System.out.println(sum);
```

### findFirst

findFirst is used to get the first element from the stream. It returns Optional.

### Example

```
List<Integer> numbers = List.of(5, 10, 15);
```

```
Optional<Integer> first = numbers.stream()
```

```
.filter(n -> n > 6)
```

```
.findFirst();
```

```
System.out.println(first.get());
```

### anyMatch

anyMatch is used to check whether any element matches a condition. It returns boolean.

### Example

```
List<Integer> numbers = List.of(3, 5, 8, 9);
```

```
boolean result = numbers.stream()
```

```
.anyMatch(n -> n % 2 == 0);
```

```
System.out.println(result);
```

## 10. Map vs FlatMap

Map and flatMap are both used to transform data in Stream API, but they are used in different situations. Map is used when one input element produces exactly one output element. It applies a function to each element and returns a stream of transformed elements. For example, if we have a list of numbers and we want to square each number, we use map because each number becomes one squared value.

FlatMap is used when one input element produces multiple output elements. It first converts each element into a stream and then flattens all those streams into a single stream. FlatMap is commonly used when working with nested collections like a list of lists. For example, if we have a list of lists of integers and want a single list of all integers, flatMap is used.

## 11. Stream vs Parallel Stream

Stream processes elements sequentially using a single thread. Each element is handled one after another in a predictable order. It is simple to understand, easy to debug, and works well for small to medium data sets or when order of execution matters. Sequential streams are generally preferred in most business applications because they are stable and easier to control.

Parallel stream processes elements concurrently using multiple threads from the common ForkJoinPool. The data is split into multiple parts and processed in parallel on different CPU cores. This can improve performance for large data sets and CPU intensive operations. However, parallel streams introduce overhead, can cause thread contention, and may lead to incorrect results if the operations are not thread safe.

## 12. What is CompletableFuture?

CompletableFuture is a class used for asynchronous and non blocking programming in Java. It allows a task to run in the background without blocking the main thread and lets us define what should happen after the task completes. Unlike traditional Future, CompletableFuture supports chaining multiple actions, handling exceptions, and combining results of multiple asynchronous tasks.

## 13. CompletableFuture vs Future

Future represents the result of an asynchronous computation, but it has several limitations. When we call get on a Future, the calling thread is blocked until the result is available. Future also does not support chaining multiple tasks or proper exception handling, which makes it difficult to build complex asynchronous flows.

CompletableFuture overcomes these limitations. It is non blocking and allows us to attach callback methods that run automatically when the task completes. We can chain multiple asynchronous operations, combine results from different tasks, and handle exceptions in a clean way. In real applications, Future is suitable for simple async tasks, while CompletableFuture is preferred for building scalable and responsive systems.

## 14. How to decide thread pool size?

For CPU bound tasks thread pool size should be close to number of CPU cores because more threads will cause context switching.

For I/O bound tasks thread pool size can be higher than number of cores because threads spend time waiting for I/O operations.

## 15. Intermediate vs Terminal operations

In Stream API, operations are divided into intermediate and terminal operations based on how they behave during stream processing. Intermediate operations are operations that return another stream and do not produce a final result. Examples are filter, map, flatMap, and sorted. These operations are lazy, meaning they do not execute immediately when they are called. They only define the processing steps and execution starts only when a terminal operation is applied.

Terminal operations are operations that produce a result or a side effect and end the stream pipeline. Examples are forEach, collect, reduce, findFirst, and anyMatch. When a terminal operation is called, the stream starts processing all intermediate operations in sequence and produces the final result. After a terminal operation, the stream cannot be reused.

## **16. How does Stream API improve performance**

Stream API improves performance mainly through lazy execution, pipelined processing, and optional parallelism. Lazy execution means that intermediate operations like filter and map are not executed immediately. They run only when a terminal operation is called, and only on the elements that are actually needed. This avoids unnecessary computations.

Streams also use pipelined processing, where each element flows through all intermediate operations before the next element is processed. This reduces the number of iterations over the data. Additionally, Stream API supports parallel streams, which can divide work across multiple CPU cores for CPU intensive tasks. Because of these reasons, Stream API often provides better performance along with cleaner and more readable code.

## **17. Difference between Optional and null and why Optional exists**

Null represents the absence of a value, but it provides no information or safety. Using null often leads to NullPointerException at runtime, which is one of the most common bugs in Java applications. With null, it is not clear whether a value can be absent or not unless clearly documented.

Optional was introduced to solve this problem. Optional is a container object that may or may not contain a value. It forces the developer to explicitly handle the absence of a value using methods like isPresent, orElse, or ifPresent. This makes the code safer and more readable. Optional exists to reduce NullPointerException, clearly express that a value may be missing, and encourage better coding practices in modern Java applications.

## **18. How does CompletableFuture improve async programming**

CompletableFuture improves asynchronous programming by removing the need to block threads while waiting for results and by providing a clean way to define what should happen after an async task completes. With older approaches like Future, the thread must call get and wait, which reduces performance and scalability. CompletableFuture allows tasks to run in the background and lets us attach callback methods that automatically execute when the result is ready.

It also supports chaining multiple asynchronous operations using methods like thenApply and thenCompose, combining results from multiple async tasks, and handling exceptions in a structured way. Because of this, CompletableFuture helps build non blocking, responsive, and scalable applications, which is especially important in modern web and microservices based systems.

## **Collection**

### **1. What is List in Java?**

List is an ordered collection that allows duplicate elements. Each element has an index, starting from zero. It is used when you want to store elements in the same order you insert them.

## **Common List Implementations**

### **ArrayList**

Stores elements in a dynamic array. Fast for reading (get) and slow for inserting in the middle. Good for most use cases.

### **LinkedList**

Stores elements as nodes linked to each other. Fast for add and remove operations in the middle. Slower for random access.

### **Vector**

Similar to ArrayList but synchronized. It is thread-safe but slower. Not used much today.

## **2. What is Set in Java?**

Set is a collection that does not allow duplicate elements. It is used when you want only unique values.

## **Common Set Implementations**

### **HashSet**

Stores elements using hash values. It does not maintain order. Very fast for add, remove, and search.

### **LinkedHashSet**

Same as HashSet but maintains insertion order. Useful when you need uniqueness + preserved order.

### **TreeSet**

Stores elements in sorted order (natural or custom). Slower than HashSet because it uses a tree structure.

## **3. What is Map in Java?**

Map stores data in key-value pairs. Keys must be unique, but values can be duplicate. It is used for fast lookups using keys.

## **Common Map Implementations**

### **HashMap**

Stores data using hashing. Does not maintain order. Very fast and most commonly used.

### **LinkedHashMap**

Maintains insertion order. Useful for caching and predictable iteration.

### **TreeMap**

Maintains keys in sorted order. Useful when you need sorted key-based operations.

### **Hashtable**

Thread-safe version of HashMap but slower. Mostly replaced by ConcurrentHashMap.

### **ConcurrentHashMap**

Thread-safe and fast. Used in multithreaded applications.

## **4. Quick Comparison Table**

List → Ordered, allows duplicates

Set → Unordered or ordered depending on implementation, no duplicates

Map → Key-value store, keys unique, values can repeat

ArrayList → Fast read, slow insert in middle

LinkedList → Fast insert/remove, slow read

HashSet → Fast, no duplicates, no order

TreeSet → Sorted set

LinkedHashSet → fast, preserves insertion order, slightly more memory.

HashMap → Fast, no order

TreeMap → Sorted map

LinkedHashMap → Maintains insertion order

## 5. ArrayList vs LinkedList

ArrayList and LinkedList are both implementations of the List interface but they are designed for different use cases. ArrayList is backed by a dynamic array, so it provides fast random access using index, which makes get operations very efficient. However, inserting or deleting elements in the middle of an ArrayList is slow because elements need to be shifted.

LinkedList is backed by a doubly linked list, so insertion and deletion operations are faster, especially in the middle of the list, because only references are updated. However, accessing an element by index is slow because it requires traversal from the beginning or end. In real applications, ArrayList is preferred in most cases because read operations are more common, while LinkedList is useful when frequent insertions and deletions are required.

## 6. What will happen if we make an ArrayList collection as final in Java?

If an ArrayList is declared final, you cannot reassign the reference to a new list, but you can still modify the list itself by adding, removing, or updating elements.

## 7. HashSet vs TreeSet:

HashSet and TreeSet are both implementations of the Set interface, but they differ in how they store and retrieve elements. HashSet stores elements using a hash table, so it does not maintain any order and provides very fast operations for add, remove, and search with average constant time performance. It allows one null element and is preferred when ordering is not required and performance is important.

TreeSet stores elements in a sorted order using a balanced tree structure, specifically a red black tree. It maintains natural ordering or custom ordering using a comparator. Operations like add and search take logarithmic time, which is slower than HashSet. TreeSet does not allow null elements because comparison is required. In real applications, HashSet is used for fast lookup, while TreeSet is used when sorted data is needed.

## 8. HashMap vs TreeMap:

HashMap and TreeMap are both implementations of the Map interface, but they are used for different purposes. HashMap stores entries using a hash table and does not maintain any order of keys. It provides very fast performance for put and get operations with average constant time complexity. HashMap allows one null key and multiple null values and is commonly used when fast access is required and ordering is not important.

TreeMap stores entries in a sorted order based on the natural ordering of keys or a custom comparator. Internally it uses a red black tree, so operations like put, get, and remove take logarithmic time. TreeMap does not allow null keys because it needs to compare keys. In real applications, HashMap is preferred for performance, while TreeMap is used when sorted keys or range based operations are required.

## 9. How to create custom ArrayList which doesn't allow duplicate?

```
public class CustomArrayList extends ArrayList {  
  
    @Override  
    public boolean add(Object o) {  
        if(this.contains(o)){  
            return true;  
        }else{  
            return super.add(o);  
        }  
    }  
}
```

## 10. Why Set doesn't allow duplicates?

Set implementations in Java internally use a Map. For example, HashSet uses a HashMap internally, and every element you insert into the Set is stored as a **key** in the map, with a constant dummy value (usually Boolean.TRUE). Since a Map cannot have duplicate keys, the Set automatically prevents duplicate elements.

TreeSet is different: it uses a TreeMap internally, where elements become keys in a sorted tree structure.

## 11. Comparable vs Comparator:

### Comparable

Comparable is used when a class defines its **own natural sorting order**. The class implements Comparable and overrides the compareTo() method. This means the sorting logic is inside the object itself. It is useful when you have a single default sorting rule, like sorting by id or name.

### Comparator

Comparator is used when you want to define **multiple or custom sorting rules** outside the class. It is implemented separately and passed to sorting methods. This allows sorting the same object in different ways, like sorting employees by name, salary, or age.

## 12. Fail -fast and Fail-safe Iterator:

A fail fast iterator immediately throws ConcurrentModificationException if the collection is structurally modified while iterating, except through the iterator's own remove method. This behavior helps detect bugs early. Iterators of collections like ArrayList, HashMap, and HashSet are fail fast. For example, if one thread iterates over an ArrayList and another thread modifies it, the iterator will fail immediately.

Fail safe iterators do not throw ConcurrentModificationException. They work on a copy of the original collection, so modifications do not affect the iterator. Because of this, fail safe iterators may not reflect the latest changes but they provide safer iteration in concurrent environments. Iterators of collections like ConcurrentHashMap and CopyOnWriteArrayList are fail safe. In simple terms, fail fast detects problems quickly, while fail safe avoids exceptions by working on a snapshot.

### 13. Why do we need ConcurrentHashMap if Hashtable is already synchronized?

Hashtable is synchronized, but it uses full table locking. This means the entire Hashtable is locked for every read and write. Because of this, only one thread can access it at a time, which makes it very slow in multithreaded applications.

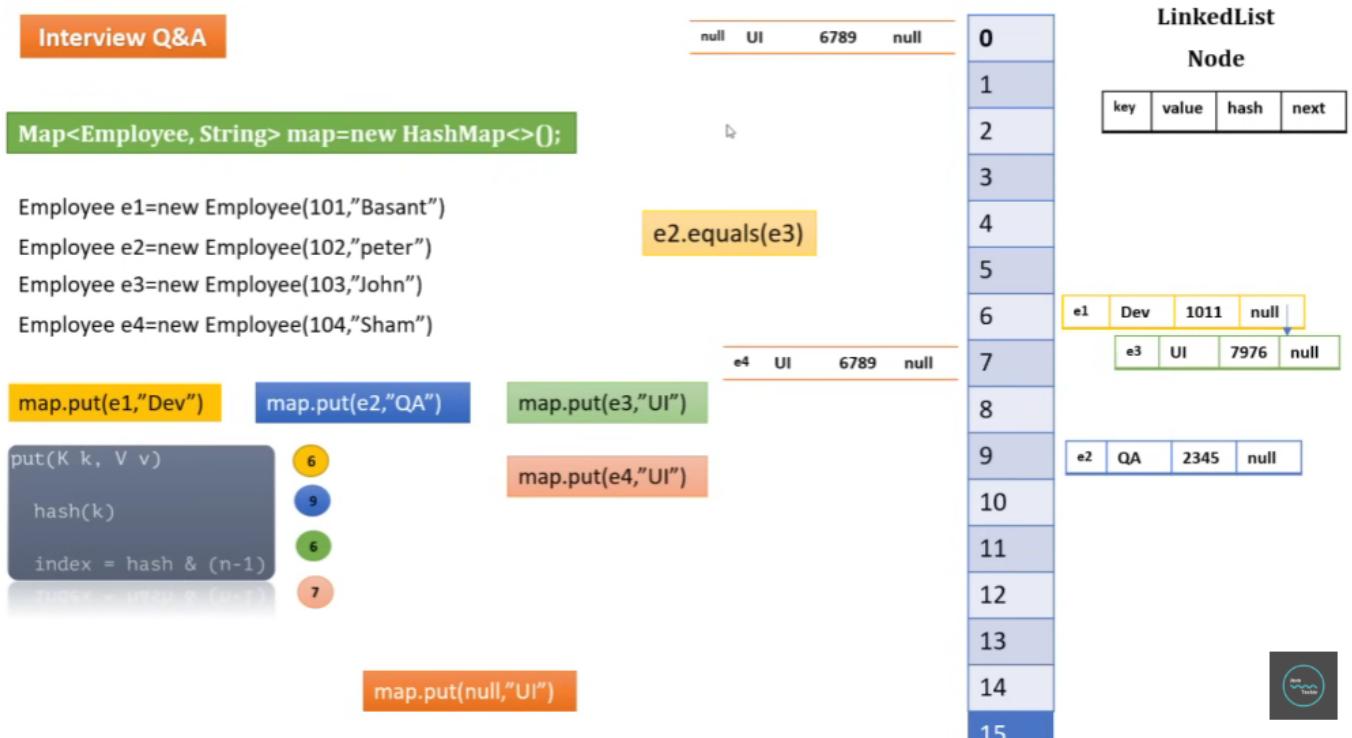
ConcurrentHashMap solves this by using partial locking (also called segment-level or bucket-level locking). Only small parts of the map are locked during updates, and reads happen almost without locking. This gives much higher performance in multi-threaded environments.

### 14. How does HashMap work internally?

Internally, HashMap uses an array of nodes, where each node contains a key, value, hash, and a reference to the next node (in case of collision). In Java 8 and above, these nodes can also form a **balanced tree** (Red-Black Tree) when collisions exceed a threshold.

When we insert a key-value pair:

- First, it calculates the hash code of the key.
- Then, it finds the index in the array using  $(n - 1) \& \text{hash}$ .
- If the bucket is empty, it stores the node directly.
- If not, it checks for collision. If the key already exists (checked via `equals()`), the value is updated. If not, it adds the new node to the end of the list or tree."



### 15. If a key is null in HashMap then where that entry will store in map?

At the 0<sup>th</sup> node.

### 16. How TreeMap Internally Works (Simple, Interview-Perfect Answer)

TreeMap stores key-value pairs in sorted order using a Red-Black Tree, which is a self-balancing binary search tree.

Every time you insert, search, or delete a key, TreeMap uses the tree structure to keep elements sorted and balanced.

Here is the step-by-step internal working:

- TreeMap stores keys in a Red-Black Tree

Each entry is stored as a Node containing:

- key
- value
- left child
- right child
- parent
- color (RED or BLACK)

Because it is sorted, your keys must be Comparable or you must provide a Comparator.

- Insertion uses Binary Search Tree logic

When you call:

```
treeMap.put(key, value);
```

TreeMap compares the key with the root:

- If new key < root → go left
- If new key > root → go right
- If key equals existing key → update value

Insertion always follows compareTo() or the provided Comparator.

## 17. How ConcurrentHashMap works internally?

- The map is divided into segments/buckets.
- When one thread updates bucket A, another thread can update bucket B without waiting.
- Read operations are non-blocking (no lock required in most cases).
- For updates, it uses CAS (Compare-And-Swap) + bucket-level locking.

**\*\*it doesn't allow duplicate keys or values because Null creates ambiguity in multi-threading ("key absent or value null?").**

## 18. How CopyOnWriteArrayList works internally?

CopyOnWriteArrayList is a thread-safe list that uses a *copy-on-write* strategy. It stores elements in a volatile array, and read operations never lock because the array is never modified in place. For every write—like add or remove—it takes a lock, creates a new copy of the entire array, applies the change, and then

replaces the old array reference. Its iterator is fail-safe because it works on a snapshot of the array. It gives very fast reads, but writes are expensive, so it's ideal for read-heavy, write-rare scenarios.

## **19 Can we use a custom class as a key in the HashMap and what is necessary for that?**

Yes, a custom class can be used as a key in a HashMap. The class must override both the hashCode() and equals() methods to ensure the correct working of the hash-based data structure. Without these overrides, it may lead to incorrect behaviour, such as failing to retrieve the correct value.

## **20. Why is HashSet faster than ArrayList for search**

HashSet is faster than ArrayList for search because of the data structure it uses internally. HashSet internally uses a HashMap to store elements. When we search for an element in a HashSet, the hashCode of the object is calculated and used to directly locate the bucket where the element should exist. Then equals is used only if needed to confirm the match. Because of this direct access, the average time complexity of search in HashSet is constant time.

ArrayList, on the other hand, stores elements in a sequential order. When searching in an ArrayList, it has to compare the target element with each element one by one using equals until it finds a match or reaches the end of the list. This results in linear time complexity and becomes slower as the list grows.

## **Multi-threading**

### **1. How does multithreading improve performance and when does it fail?**

Multithreading improves performance by allowing multiple tasks to run at the same time using different CPU cores. It is especially useful for IO bound tasks like database calls or network requests because while one thread is waiting, another thread can do useful work. However, multithreading fails when there are too many threads, which causes excessive context switching, shared resource contention, and synchronization overhead. In such cases, performance can actually become worse instead of better.

### **2. Difference between synchronized method and synchronized block**

A synchronized method locks the entire object for the whole duration of the method execution. This means no other thread can access any synchronized method of that object at the same time. A synchronized block locks only a specific section of code or a specific object. Synchronized blocks are preferred because they reduce the scope of locking and improve performance by allowing more concurrency.

### **3. What is race condition? Real production example**

A race condition occurs when multiple threads access and modify shared data at the same time, and the final result depends on the order of execution. A real production example is two threads updating the same bank account balance simultaneously. If both threads read the same balance and update it without synchronization, the final balance can be incorrect. Race conditions are prevented using synchronization, locks, or atomic variables.

### **4. What is volatile keyword and where is it used?**

The volatile keyword ensures visibility of changes across threads. When a variable is declared volatile, any update made by one thread is immediately visible to other threads. It does not provide atomicity, so it cannot be used for compound operations like increment. Volatile is commonly used for flags or status variables, such as stopping a background thread safely.

## **5. Deadlock what is it and how do you prevent it?**

Deadlock occurs when two or more threads are waiting for each other's locks indefinitely, so none of them can proceed. For example, Thread A holds Lock 1 and waits for Lock 2, while Thread B holds Lock 2 and waits for Lock 1. Deadlocks can be prevented by acquiring locks in a consistent order, avoiding nested locks, using timeouts, or using higher level concurrency utilities.

## **6. Difference between Runnable and Callable**

Runnable represents a task that does not return any result and cannot throw checked exceptions. Callable represents a task that returns a result and can throw checked exceptions. Callable is used with ExecutorService and Future to get the result of asynchronous computation. In real applications, Callable is preferred when the task outcome is required.

## **7. Explain ExecutorService and thread pool internals**

ExecutorService is a framework that manages a pool of reusable threads. Instead of creating a new thread for every task, tasks are submitted to the ExecutorService and executed by available threads from the pool. Internally, it uses a work queue to store tasks and worker threads to execute them. This improves performance, controls resource usage, and provides lifecycle management through shutdown methods.

## **8. How does Java handle thread safety in collections?**

Java provides synchronized collections and concurrent collections. Synchronized collections lock the entire collection for each operation, which reduces performance. Concurrent collections like ConcurrentHashMap use fine grained locking or lock free techniques, allowing multiple threads to read and write concurrently with better scalability and performance.

## **9. How do you decide thread pool size for CPU-bound vs IO-bound tasks?**

For CPU bound tasks, thread pool size should be close to the number of CPU cores because more threads cause unnecessary context switching. For IO bound tasks, thread pool size can be higher than the number of cores because threads spend time waiting for IO operations. The idea is to keep the CPU busy while other threads are waiting.

## **10. Difference between wait, notify, and notifyAll**

wait causes the current thread to release the lock and go into waiting state until it is notified. notify wakes up one waiting thread, while notifyAll wakes up all waiting threads. These methods are used for inter thread communication and must be called inside a synchronized block or method.

## **11. What is ThreadLocal and where is it used?**

ThreadLocal provides thread level storage, meaning each thread has its own independent copy of a variable. This avoids synchronization because data is not shared between threads. ThreadLocal is commonly used to store user context, transaction information, or request specific data in web applications and frameworks like Spring.

## **SCENARIO-BASED:**

## **12. A thread is stuck in waiting state. How do you identify and fix it?**

If a thread is stuck in waiting state, the first thing I do is take a thread dump using tools like jstack or JVM monitoring tools. This shows which thread is waiting and on which lock or condition. Usually, the issue is a missing notify or notifyAll call, or improper use of wait inside synchronized blocks. To fix it, I ensure wait notify are used correctly, conditions are rechecked in a loop, and locks are properly released so the waiting thread can resume.

### **13. Two threads update the same bank account balance. How do you prevent inconsistency?**

This is a classic race condition. To prevent inconsistency, I ensure that only one thread can update the balance at a time. This can be done using synchronized blocks, ReentrantLock, or atomic variables if the logic is simple. In real applications, account updates are usually wrapped inside a transactional or locked section so read modify write happens atomically.

### **14. Application slows down due to shared resource contention. Your approach?**

First, I identify which resource is causing contention using thread dumps or profiling tools. Then I reduce the lock scope by synchronizing only the critical section instead of whole methods. I may also use ReadWriteLock, concurrent collections, caching, or split the shared resource to reduce contention. The goal is to minimize blocking between threads.

### **15. How do you safely stop a long-running background thread?**

Threads should not be stopped forcefully. The safe way is to use thread interruption. I set an interrupt flag using interrupt and inside the thread logic I regularly check Thread.interrupted and exit gracefully. This allows cleanup of resources and prevents inconsistent state.

### **16. How do you schedule a task every 5 seconds?**

I use ScheduledExecutorService. It provides methods like scheduleAtFixedRate or scheduleWithFixedDelay. This approach is better than using Timer because it handles exceptions properly and supports thread pooling, which makes it suitable for production systems.

### **17. Two threads must execute in strict order A then B. How?**

To ensure execution order, I use coordination mechanisms like CountDownLatch or wait notify. For example, Thread B waits on a latch until Thread A finishes its work and counts down. This guarantees ordering without busy waiting and is safe and readable.

### **18. ExecutorService rejecting tasks. Possible reasons?**

Tasks are rejected when the thread pool is exhausted. Common reasons are all threads are busy, the task queue is full, or the maximum pool size is reached. Another reason is submitting tasks after shutdown. The rejection policy then decides how tasks are handled. Fix involves tuning pool size or queue capacity.

### **19. Shared counter gives random values. What's happening?**

This is due to a race condition. Increment operation is not atomic and multiple threads update the counter at the same time. To fix this, I use AtomicInteger or synchronize the increment operation so updates are thread safe.

### **20. Need to wait for multiple threads to finish. What will you use?**

I use CountDownLatch or ExecutorService with awaitTermination. CountDownLatch allows the main thread to wait until all worker threads complete. This is clean and commonly used in concurrent workflows.

### **21. Thread blocked on IO needs cancellation. What will you do?**

For IO blocking, interrupt alone may not work. I close the underlying IO resource like socket or stream, which causes the blocked thread to exit. If possible, I use interruptible IO like NIO channels that respond to thread interruption.

### **22. Works locally but fails in production. What do you check first?**

I first check environment differences such as CPU cores, memory limits, thread pool sizes, and load. Then I look for race conditions or missing synchronization that appear only under high concurrency. Logs and thread dumps from production help identify the root cause.

### **23. CPU-heavy task. More threads or fewer threads and why?**

For CPU heavy tasks, fewer threads are better. Ideally, thread count should be close to the number of CPU cores. More threads cause context switching overhead and reduce performance instead of improving it.

### **24. Frequent deadlocks. How do you detect and eliminate them?**

Deadlocks are detected using thread dumps or JVM tools that show circular lock dependencies. To eliminate them, I enforce consistent lock ordering, reduce nested locking, or use tryLock with timeout so threads can recover instead of waiting forever.

### **25. ExecutorService not shutting down after shutdown. Fix?**

Shutdown only stops accepting new tasks but allows running tasks to finish. If it does not stop, I call awaitTermination with a timeout and then shutdownNow to interrupt running tasks. I also check for blocked threads or infinite loops inside tasks.

### **26. Share data between threads without synchronized. Options?**

I use volatile variables for visibility, atomic classes for atomic operations, concurrent collections like ConcurrentHashMap, or ThreadLocal when data should be isolated per thread. These options reduce locking and improve performance while maintaining thread safety.

### **27. What problem does the Executor framework solve compared to creating threads manually?**

The Executor framework solves the problem of uncontrolled thread creation and poor resource management. When we create threads manually, each task creates a new thread, which is expensive and can exhaust CPU and memory. Executor framework separates task submission from task execution. We submit tasks, and the framework manages thread creation, reuse, scheduling, and lifecycle. This results in better performance, scalability, and easier maintenance in production systems.

### **28. What is Executor, ExecutorService, and ScheduledExecutorService?**

Executor is a simple interface that executes a task without managing thread lifecycle. ExecutorService extends Executor and adds features like task submission, Future support, and shutdown control. ScheduledExecutorService extends ExecutorService and is used to run tasks after a delay or periodically. In real applications, ExecutorService and ScheduledExecutorService are commonly used for async and scheduled tasks.

### **29. How does ExecutorService internally manage threads and tasks?**

ExecutorService internally uses a thread pool, a task queue, and worker threads. When a task is submitted, it is placed into a queue. If a thread is available, it picks up the task and executes it. If not, the task waits in the queue. Threads are reused instead of being destroyed, which reduces overhead and improves performance.

### **30. Difference between execute and submit?**

execute is used to run a Runnable task and does not return any result or Future. If an exception occurs, it is thrown directly. submit can be used with Runnable or Callable and returns a Future object. Using Future, we can check task status, get results, or handle exceptions. submit is preferred when task monitoring or result is required.

### **31. What is a thread pool and why is it better than creating new threads?**

A thread pool is a group of reusable threads managed by ExecutorService. Instead of creating a new thread for every task, tasks are executed by existing threads. This reduces thread creation overhead, improves

performance, controls resource usage, and prevents system overload. Thread pools are essential in high concurrency applications.

### **32. Difference between newFixedThreadPool and newCachedThreadPool?**

newFixedThreadPool creates a fixed number of threads and uses a queue to store extra tasks. It provides predictable resource usage. newCachedThreadPool creates threads as needed and reuses idle threads, but it has no upper limit. Cached thread pool can lead to resource exhaustion if tasks increase rapidly, so fixed thread pool is safer for production.

### **33. What happens internally when a task is submitted to ExecutorService?**

When a task is submitted, ExecutorService checks if an idle thread is available. If yes, the task is assigned to that thread. If not, the task is placed in the queue. If the queue is full and thread limit allows, a new thread is created. Otherwise, the task is rejected based on the rejection policy. This controlled flow ensures stability.

### **34. How does Future work with ExecutorService?**

Future represents the result of an asynchronous task. When a task is submitted using submit, a Future is returned. Using Future, we can check whether the task is completed, cancel it, or retrieve the result using get. This allows non blocking task submission and controlled result handling.

### **35. Difference between Runnable and Callable in Executor framework?**

Runnable does not return a result and cannot throw checked exceptions. Callable returns a result and can throw checked exceptions. Callable is preferred when the task outcome is required or exception handling is important. Both are executed by ExecutorService.

### **36. How do you properly shut down an ExecutorService?**

To properly shut down ExecutorService, first call shutdown to stop accepting new tasks and allow running tasks to finish. Then use awaitTermination to wait for completion. If tasks do not finish, call shutdownNow to interrupt running threads. This ensures graceful shutdown and avoids resource leaks.

### **37. Can you change thread pool size at runtime?**

Yes, thread pool size can be changed at runtime, but only if we are using ThreadPoolExecutor directly. ExecutorService created using Executors utility methods does not expose methods to change the size. With ThreadPoolExecutor, we can change corePoolSize and maximumPoolSize using setter methods while the application is running. This is useful in production when load changes, for example increasing threads during peak traffic and reducing them later. However, changing pool size should be done carefully because increasing threads blindly can cause CPU contention and memory issues. In real systems, this is usually combined with monitoring and tuning rather than frequent dynamic changes.

### **38. Why should threads be reused instead of recreated?**

Threads are expensive to create and destroy because they consume memory and involve interaction with the operating system. Creating a new thread for every task adds overhead and slows down the application under load. Reusing threads through a thread pool avoids this cost by keeping threads alive and assigning them new tasks when needed. This improves performance, reduces latency, and gives better control over resource usage. In production systems, thread reuse also prevents system overload and makes application behaviour more predictable under high concurrency.

## Miscellaneous

### **1. I have the same method inside an abstract class and an interface. I implement both in a class. Which method gets invoked and why?**

In this situation, the method from the abstract class is always invoked, not the interface method. This is because in Java, class methods have higher priority than interface default **methods** during method resolution.

When a class extends an abstract class and also implements an interface, Java follows a clear rule called class wins over interface. Even if the interface has a default method with the same signature, the JVM will always choose the method from the abstract class. The reason is that class inheritance existed before default methods were introduced in Java 8, so Java gives priority to class hierarchy to avoid ambiguity and to maintain backward compatibility.

For example, if an abstract class defines a concrete method and an interface defines a default method with the same name and parameters, and a class extends the abstract class and implements the interface, then calling that method on the object will execute the abstract class's method. The interface default method is ignored unless the class explicitly overrides the method and chooses to call the interface version using `InterfaceName.super.methodName`, which is optional.